

RSLAB

A Pure-Rust Sparse Direct Solver: Algorithms, Implementation, and Benchmarks

Milan Rother

August 31, 2026

Abstract

RSLAB is a pure-Rust sparse direct solver with three factorization paths: a supernodal complex-symmetric $\mathbf{LDL}^\top$ with Bunch–Kaufman pivoting, an unsymmetric supernodal/multifrontal $\mathbf{LU}$ with threshold partial pivoting, and a KLU-style path (block triangular form + per-block Gilbert–Peierls $\mathbf{LU}$) for circuit-shaped systems. It is generic over the scalar type (`f64`, `f32`, and their complex counterparts) and links no external numeric library: the SIMD dense kernels, the fill-reducing orderings, and the factorizations are all Rust. The analyze phase computes a-priori peak-memory, work, and runtime predictors as exact functions of the matrix structure; the default configuration is deterministic and model-free (an exact ordering race that scores candidates by their true factor fill, the measured kernel knobs, and a worker count from a one-time cached hardware calibration); and the numeric factorization is bit-identical regardless of the worker count. A single-precision factor can precondition a double-precision Krylov solve, recovering the accuracy at half the factor memory. When the factor is made approximate (static-pivoted or thresholded) it becomes a preconditioner, and a matching Krylov layer (COCG/COCR, flexible GMRES, batched multi-RHS block GMRES with within-cycle deflation, and GCRO-DR subspace recycling) closes the accuracy gap. This report documents the core algorithms, their implementation, and the measured performance against `faer`, `MKL`, `PARDISO` and Apple Accelerate over a complete-distribution corpus.

1 Scope and architecture

RSLAB factorizes and solves $Ax = b$ for sparse A that is complex-symmetric (routed to $\mathbf{LDL}^\top$), general unsymmetric (routed to $\mathbf{LU}$), or circuit-shaped (routed to the KLU path). The primary target is the complex-symmetric systems of frequency-domain electromagnetic and FEM analysis (curl–curl Maxwell problems with lossy media or PML are complex-symmetric rather than Hermitian), with the KLU path serving the MNA/SPICE-class operators of circuit extraction. The implementation is stable Rust with no C/Fortran dependency and no BLAS/LAPACK/METIS linkage, and is exposed to Python through a thin PyO3 layer that mirrors the Rust surface (`rslab.ldlt/lu/klu/spsolve`).

All paths follow the standard three-phase flow. *Analyze* orders the matrix, builds the elimination structure, and produces the a-priori predictors; *factor* performs the numeric decomposition on the ordered, equilibrated matrix; *solve* runs the triangular substitutions. Analyze is separated from factor at the API level: a symbolic object (`LdltSymbolic`, `LuSymbolic`, `KluSymbolic`) is produced once and factors any number of matrices sharing its sparsity pattern, so a frequency sweep or Newton iteration pays the ordering and symbolic cost once. Across the corpus the numeric factor dominates the wall-clock, so the tuning and the parallelism target the factor. The factored solvers implement the crate’s `Preconditioner` and `Factorization` traits, which is how the direct and iterative layers compose.

2 Fill-reducing orderings

The ordering determines fill and therefore both factor time and memory. RSLAB implements three ordering families behind an `OrderingMethod` enum and an adaptive dispatcher, sharing one quotient-graph elimination engine.

Approximate minimum degree. AMD [1] selects the pivot of least approximate external degree from degree-indexed bucket lists over a sliding index workspace with in-place garbage collection. The update performs aggressive element absorption, mass elimination, and supervariable detection by adjacency hashing; rows of degree above $\max(16, \lfloor 10\sqrt{n} \rfloor)$ are deferred to the tail as dense. AMD is the `Auto` pick for very large, very sparse patterns, and a candidate in the race below.

Approximate minimum fill. AMF (HAMF4 [2, 3]) replaces the degree score with an approximate fill score $\text{RMF}(i) = (\deg(i)(\deg(i) - 1 + 2\deg_{me}) - \text{WF}(i)) / (nv(i) + 1)$, where the working fill `WF` is accumulated per eliminated element in 64-bit integers. It is the `Auto` pick for $n \leq 10\,000$ and, in practice, the race’s most frequent winner on grid problems, staying within $1.10\times$ of the reference HAMF4 fill on a 183 277-matrix oracle suite.

Nested dissection. A multilevel recursive bisection [4, 5] driven by an explicit work stack (no recursion): coarsen, bisect, refine with Fiduccia–Mattheyses, uncoarsen with projection and refinement. The separator refinement is a clean-room node-separator scheme (a two-sided node-FM pass over the separator’s boundary with balance-aware gain buckets, projected and re-refined at every level) rather than a port; on a 40^3 grid it produces 11.24M factor entries against 12.52M for MKL METIS. Subgraphs below 200 vertices or splits with a $\geq 90\%$ side fall back to AMD leaves. The bisection seed is itself raced: a small ensemble of seeds is scored by exact fill and the best kept, deterministically.

Dispatch. `OrderingMethod::Auto` routes by cheap structural predicates (very large and sparse $\rightarrow$ AMD; arrow/bordered patterns $\rightarrow$ AMF; otherwise a size rule). The shipped default is `AutoRace`, a two-stage *exact* race. Stage one runs the {AMD, AMF, RCM} symbolic prefixes concurrently and keeps the smallest true factor nnz; stage two admits the nested-dissection candidate only above a work floor (5×10^9 predicted flops and $n > 10\,000$), where its analysis cost amortizes against the factor it saves, and adopts it only on a strictly smaller exact fill. Ties break deterministically, so the pick is a pure function of the pattern. The race costs about $4\times$ a single-pass analysis and returns the exact minimum over its candidates rather than a predicted one.

3 Symbolic analysis

From the ordered pattern RSLAB builds the elimination tree and detects fundamental supernodes, maximal column runs whose row structure differs only by the eliminated column. Small supernodes are amalgamated to widen the dense fronts that feed the BLAS-3 kernels: nodes with fewer than `nemin` = 16 columns [6] merge with their parent, using a column renumbering that re-postorders the tree so a desired parent–child merge becomes structurally adjacent. Column counts from a depth-first tree walk drive the factor-size estimate.

Relaxed amalgamation [7] merges an adjacent child within a width and extra-row budget, padding the merged front with explicit zeros to trade fill for higher-rank Schur updates. It is available (`with_relax`) and off by default. The padded rows are structurally zero, but the numeric path carries them through every update as arithmetic and as memory traffic, so the trade is favourable only when the padding is small relative to a front that is already wide. On grid operators, whose fundamental supernodes are narrow and numerous, it is not: a 145^2 convection–diffusion grid has 3878 fundamental supernodes against 277 relaxed ones at identical output fill, and factors $3.4\times$ slower with the relaxation on. Over the 18-matrix corpus the default is 0.654 in geomean factor time against relaxed merging, at identical or lower fill.

4 Numeric factorization

4.1 Equilibration

Before factoring, A is symmetrically equilibrated as $\hat{A} = DAD$ with $s_i = 1/\sqrt{\max_j |A_{ij}|}$ in one pass over the lower triangle; the real, symmetric, positive diagonal preserves complex symmetry and the pivot signs and tolerates zero diagonals (saddle-point systems). The **LU** path applies a two-sided $\hat{A} = D_r A D_c$; iterative Ruiz and MC64 matching-based scaling are selectable alternatives.

4.2 Supernodal left-looking LDL[⊤]

The default symmetric path assembles one dense panel per supernode, updates it by every previously factored descendant (`cmod`: pull the descendant’s columns, apply its block-diagonal D , subtract the rank update, a scalar triple loop below a flop threshold, a SIMD GEMM above), then factors the fully-summed block in place (`cddiv`) by a blocked Bunch–Kaufman partial factorization [8] with panel width 64. For complex-symmetric A the kernel is the real `sytf2` control flow with magnitudes $|z|$ and no conjugation; the pivot threshold is $\alpha = (1 + \sqrt{17})/8$, pivot selection is bounded to the supernode’s fully-summed rows, and a 2×2 block is rejected as singular when $|d_{11}d_{22} - d_{21}^2| < \varepsilon \max(|d_{11}|, |d_{21}|, |d_{22}|)^2$ with $\varepsilon = 10^{-14}$, bounding the element growth injected into the trailing update.

Transient memory is controlled by *panel freeing*: every supernode carries a reference count of the ancestors that will still update from it, and when a `cmod` decrements it to zero the panel is compacted into its final CSC fragment and the dense buffer released (acquire/release atomics make this safe under tree parallelism). The resident data is the compact factor plus only the live panels. Sibling subtrees factor concurrently in a scoped rayon pool of caller-set width; per-node GEMMs parallelize above their flop thresholds.

4.3 Multifrontal path

The alternative multifrontal schedule [9, 10] assembles a dense frontal matrix per supernode from its eliminated columns and its children’s contribution blocks (CBs), factors the fully-summed part, and passes the Schur complement up as its CB. A work-stealing tree recursion factors children in parallel; each CB is freed the moment it is extend-added into the parent, and where the CB stack is large children are reordered by Liu’s (peak – *cb*) rule [11], a scheduling hint that changes neither numbering nor fill. The front is denser and more BLAS-3-rich than a left-looking panel but carries a higher transient peak, which is why both schedules ship and the choice is per matrix.

4.4 Unsymmetric LU

General matrices reuse the symmetric analysis machinery on the symmetrized pattern $A \cup A^{\top}$ and ship both numeric schedules. The default supernodal left-looking **LU** uses UMFPACK-style [12] threshold partial pivoting: the diagonal candidate is kept when $|a_{jj}| \geq u |a_{\text{colmax}}|$ with tunable $u = \text{pivot_u}$ (default 0.1); $u = 1$ recovers full partial pivoting, $u = 0$ keeps the natural order and skips the pivot search (the static fast path for fixed-pattern sequences). The multifrontal **LU** factors each front by a blocked right-looking `getrf` with unrestricted partial pivoting over the fully-summed rows. L is stored CSC (unit lower), U CSR; the column permutation is the fill-reducing ordering and a separate row permutation records the interchanges.

4.5 KLU path for circuit-shaped matrices

Circuit-class operators (MNA/SPICE: extremely sparse, unsymmetric, near-triangularizable) defeat supernodal methods: their irreducible blocks are far too small for BLAS-3 fronts to pay off. RSLAB therefore ships a third path following SuiteSparse KLU [14], reimplemented in pure Rust.

Analyze permutes to block *upper* triangular form (BTF). A maximum transversal pairs every column with a row holding a structural nonzero (an incomplete matching proves the matrix *structurally* singular before any numeric work), and Tarjan’s SCC algorithm on the matched digraph, emitted in reverse topological order, yields the symmetric permutation under which every entry outside the irreducible diagonal blocks lies above them. Two transversal algorithms ship, the MC21-style augmenting-path matching with per-column lookahead and Hopcroft–Karp [16]; a maximum matching is not unique, and the choice changes the block structure. Both candidates are therefore scored by the pattern-only Gilbert–Peierls fill count of the blocks they induce, and the smaller wins. A diagonal-repair pass then routes the matching through off-diagonal pivots where that keeps the diagonal structurally nonzero, which removes a fill cascade seen under row scaling. AMD orders each block on its symmetrized pattern. Fill is confined to the diagonal blocks; the off-block entries are never factored and enter only the block back-substitution.

Factor runs a left-looking Gilbert–Peierls **LU** per block: for each column a depth-first reach over the pattern of the L columns factored so far yields the fill pattern in topological order, so the numeric work is proportional to the flop count. Threshold partial pivoting prefers the (structurally nonzero) diagonal unless it falls below 10^{-3} of the column maximum; row-max scaling equilibrates the rows first. Results are **bit-identical across runs and thread counts** in every mode, which is what lets the path double as the determinism arbiter for the supernodal ones.

Two levels of parallelism sit behind that guarantee, both gated by one principle: a *work floor* (the overhead is real and is never bypassed) and a *concurrency ratio* of at least two (an exact Amdahl bound across blocks, the mean elimination-DAG level width within a block, which also bounds the worker count). Above the gate, the independent BTF diagonal blocks factor concurrently with a two-phase spliced output. Refactorization additionally pipelines the columns *inside* a dominant irreducible block in the manner of NICSLU [17]: worker-owned column ranges on the frozen elimination DAG, with just-in-time spin-waits on per-column ready flags instead of level barriers, on OS threads rather than a work-stealing pool (a spinning task must not starve the column owner it waits for). That is worth 2.0 to 2.9 $\times$ on the work-heavy SuiteSparse circuit matrices (rajat15, twotone, ASIC_100ks, onetone1) and changes no bit of the result.

Refactor is the path’s distinctive fast lane: for a new value set on the *same* pattern (frequency sweep, time stepping, Newton), the stored pattern and pivot sequence are replayed with no symbolic work and no pivot search; a changed pattern is detected and rejected exactly, and a frozen pivot that becomes zero fails cleanly so the caller can re-factor with pivoting. The symbolic object also carries an exact a-priori fill/flop estimator (a pattern-only Gilbert–Peierls pass under the diagonal-pivot assumption, computed lazily and cached), mirroring the estimators of the other two paths.

4.6 Static pivoting and preconditioner mode

The ZeroPivotAction policy governs near-singular pivots on the $\mathbf{LDL}^\top/\mathbf{LU}$ paths: **Fail** (default) returns a rank-deficiency error, **PerturbToEps** lifts any sub-floor pivot to the floor while preserving its phase, so the factorization satisfies $\mathbf{LDL}^\top = A + \Delta$ with bounded Δ and never fails. Together with threshold dropping (**drop_tol**: incomplete factors) and **f32** mixed precision, this never-fail factor is the preconditioner the Krylov layer consumes. The KLU path never perturbs: a vanishing pivot is a hard, attributed error.

5 A-priori predictors

The analyze phase computes a set of predictors as *exact functions of the symbolic structure* and the scalar size v , before any numeric work, the basis of the memory guarantee, the heuristic configuration pick, the thread policy, and the resource planner of Sec. 7.

Per-path peak memory. The factor size is `factor_nnz` $\cdot (v + 8)$, a structural upper bound (cancellation can only lower it). The transient peak is bounded separately per schedule. For the

left-looking path, a reference-count simulation of the panel-freeing schedule walks the postorder, adds each panel as it is assembled, frees it when its refcount reaches zero, and records the running maximum of live panels plus already-compacted factor, the floor the solver actually operates at; a conservative whole-run bound adds all panels, the input, and a scratch term $\frac{1}{4}(\text{panels} + \text{factor}) + 32 \text{ MB}$ for the per-thread `cmod/cdiv` buffers. For the multifrontal path a level-parallel model takes, per assembly-tree level, the front sizes $\sum \text{row}^2 v$ plus that level's child contribution blocks, times $\frac{7}{5}$ for work-stealing overlap. Over the corpus both bounds hold at an estimate/measured ratio of ≈ 1.3 in geomean (Fig. 1); the panel-freeing floor is the tighter quantity the resource planner of Sec. 7 budgets against. The KLU path carries the same contract: a pattern-only Gilbert–Peierls pass (Sec. 4.5) gives its fill and flops exactly under diagonal pivoting.




Figure 1: A-priori memory estimate (conservative bound and panel-freeing floor) vs. the measured peak, per RSLAB path. The estimate does not under-predict.

Work, runtime, and threads. The geometric work factor $\text{flops} = \sum_s \text{row}_s^2 \text{ncol}_s$ is a type-independent proxy; the thread-aware runtime estimate is

$$\max\left(\frac{\text{critical_path_flops}}{r}, \frac{\text{factor_flops}}{r \cdot \text{speedup}}\right), \quad r = \text{gflops} \cdot 10^9,$$

where the first term is an Amdahl floor from the assembly tree’s longest leaf-to-root chain, so a critical-path-bound matrix does not benefit from more workers. The single-thread rate and parallel speedup come from a one-time hardware calibration cached by machine fingerprint; a learned additive residual on the speedup curve (ridge regression on dimensionless structure features, dominated by `critical_path/factor_flops`) cuts the held-out speedup error by $\sim 26\%$. A thread-count predictor caps thin/narrow problems at two workers and small ones at four, landing within $\sim 10\%$ of the per-matrix optimum in geomean (a fixed two-worker budget is off by $\sim 50\%$).

6 Determinism and the shipped default

Three determinism properties hold. The numeric factor is bit-identical regardless of worker count: parallel subtrees and GEMMs write disjoint blocks and every reduction sums in a fixed order, so the thread count is a performance parameter only (the parallel multi-RHS solve is likewise bit-identical to serial). The analyze-phase choices are deterministic functions of the pattern. And the resource planner is a pure function of its inputs. The pivot sequence is value-dependent; `pivot_u = 0` (or the KLU refactor) freezes it across same-pattern refactorizations.

The shipped default configuration is deterministic and model-free: the exact two-stage ordering race of Sec. 2, the measured default kernel knobs, and, after a one-time cached hardware calibration, a critical-path-aware worker count from the calibrated cost model. Every input to the pick is an exact a-priori quantity; nothing is modeled, and nothing is measured implicitly at solve time.

Mixed precision enters through the preconditioner rather than through a separate solver: an `f32` factor at half the memory (`LowPrecisionPreconditioner/LowPrecisionLu`) preconditions an `f64` Krylov solve, GMRES-IR in the sense of Carson–Higham [15] with the working precision fixed at `f64`. On the reference class the single-precision factor runs $1.64\times$ the double one.

7 Solver-in-the-loop and resource governance

RSLAB is built to run *inside* an outer loop (a frequency sweep, adaptive mesh refinement, a Newton or time-stepping iteration, many concurrent extractions) rather than as a one-shot black box, and three mechanisms carry that model. First, analyze and factor are separate phases: a symbolic object factors any number of matrices sharing its pattern, so a same-pattern sequence pays the ordering and symbolic cost once (and on the KLU path the numeric refactor also skips the pivot search). Second, every factorization runs in a *scoped* thread pool of caller-set width, so many concurrent solves coexist without oversubscribing the machine, and a factored preconditioner hands its thread policy to the Krylov layer so factor and solve share one concurrency budget. Third, the resource planner `tuning::plan` turns a memory estimate, a budget (a memory ceiling and a thread budget), and the cached calibration into a plan (a predicted peak, a predicted runtime, and a fits/does-not-fit decision) as a pure function of its inputs. Because the memory estimate of Sec. 5 is an upper bound rather than a sample, a scheduler can pack concurrent solves against a memory ceiling without extra safety margin, fail fast before allocating, or switch to an approximate factor when the exact one cannot fit; and because per-solve peaks and runtimes are predicted, the aggregate footprint and makespan of a batch are computable before the batch runs. Per-call diagnostics (stage timings, fill, the attached estimate) close the loop after the fact, estimate against measured.

8 Preconditioned and recycled Krylov layer

When the factor is approximate, accuracy is recovered by iteration. The iteration and the preconditioner are decoupled by two traits: `LinearOperator` supplies $y \leftarrow Ax$ and a block $Y \leftarrow AX$ (an explicit matrix, or matrix-free; an FMM/MoM operator with RSLAB factoring only the sparse near-field), and `Preconditioner` supplies $z \leftarrow M^{-1}r$ and a block apply that a factored solver implements through the parallel `solve_many` (a 8 to $19\times$ wide-RHS speedup the block Krylov methods inherit).

Complex-symmetric short recurrences. For $A = A^\top$ the methods of choice are COCG [28] and COCR [29]: structurally CG/CR but with every inner product in the unconjugated bilinear form $x^\top y$. A zero bilinear denominator is reported as a deterministic breakdown with the best iterate, never a division by zero.

Flexible GMRES. For general operators, restarted GMRES [21] in the flexible variant [22]:

the preconditioned Arnoldi vectors $z_j = M^{-1}v_j$ are kept as a second basis and $x += Zy$, which saves one M^{-1} solve per restart cycle and admits a varying preconditioner. Orthogonalization is MGS with a conditional DGKS [23] second pass. Each cycle measures the *true* residual $\|b - Ax\|$, the only reliable stop test on ill-conditioned near-field operators, and a Hessenberg-diagonal guard truncates the least-squares solve to its well-conditioned prefix, so a rank-deficient cycle reports non-convergence instead of injecting NaN.

Block GMRES. The multi-RHS solver drives s right-hand sides in lockstep so matvec and preconditioner are issued once per step as block applies. Each right-hand side keeps its own Arnoldi basis (batched-independent; a shared-subspace block Krylov method is future work), which lets a column *deflate* out the moment it converges including mid-cycle, compacting the batched applies to the still-active width. The panel orthogonalization is block CGS2 [24, 25]: two panel-wide sweeps, BLAS-3 intensity, parallelized in fixed row chunks and therefore bit-identical across thread counts. Measured on preconditioned convection–diffusion ($n = 40\,000$, complex), BCGS2 scales to $\sim 2.2\times$ at 12 cores where the per-RHS MGS path is flat-to-negative, and per-RHS cost stays near-flat in the block width.

GCRO-DR subspace recycling. For sequences of related solves, GCRO-DR [26, 27] retains the k harmonic-Ritz vectors of smallest value, the near-invariant subspace that governs restarted-GMRES stagnation, and deflates them across restarts within a solve, and recycles them across solves through an opaque handle refreshed in place. The subspace only *shapes* the Krylov space: the outer true-residual test stays authoritative, every recycle step degrades gracefully to plain FGMRES, and $C = \tilde{A}U$ is recomputed each solve so the split stays exact when A or M drift. On a stagnation-spectrum sequence of 8 solves ($n = 20\,000$), recycling cuts the cumulative iteration count $6.4\times$ versus cold starts ($2.9\times$ on the first solve alone) for a $1.5\times$ wall-clock reduction. Warm starts (x_0), an adaptive restart length that keeps the up-front Arnoldi basis under a memory budget, and f32 mixed-precision preconditioning (`LowPrecisionPreconditioner` / `LowPrecisionLu`) complete the layer.

9 Benchmarks

All cross-solver figures come from the `bench_suite` engine over a complete-distribution corpus: structured-grid generators (curl–curl Maxwell, shifted Helmholtz, Stokes/KKT saddle-point, convection–diffusion over the grid-Péclet range, BEM/MoM near-field kernels) plus the complex SuiteSparse [20] matrices, 8k–125k DOFs, `Complex<f64>`, measured in one run on a quiet 12-core machine, so the cross-solver ratios carry no run-to-run drift. RSLAB runs its shipped default (the exact ordering race and the calibrated worker count); each path is compared on its own class against its own PARDISO mtype [18] and faer [19]. This corpus was measured with RSLAB v0.27 on a 12-core x86 machine; the Apple measurement of Sec. 9.2 carries the current version.

9.1 Scaling and accuracy

Figures 2 and 3 plot factor time and peak memory against nnz with per-solver power-law fits; Table 1 gives the head-to-head geomean ratios. RSLAB sits between the two on both paths: faster than the pure-Rust faer, moderately behind the hand-optimized MKL PARDISO ($5.6\times$ sym / $5.1\times$ unsym). The heuristic pick’s win over the fixed default is visible as a gap that widens with problem size (largest on the big matrices where a mispicked ordering costs most) and on the symmetric path never comes at a memory cost. On accuracy (Fig. 4), 24 of 31 corpus matrices factor to a relative residual below 10^{-8} , matching PARDISO and ahead of faer; the exact $\mathbf{LDL}^\top$ declines a handful of indefinite saddle-point matrices rather than returning a degraded solution. The determinism and equivalence properties were validated over 180 SuiteSparse matrices: right-looking vs. multifrontal, both emit modes, parallel vs. serial front subtraction, and the 32-bit compressed factor are all bit-identical.




Figure 2: $\text{LDL}^\top$ path (symmetric, PARDISO mtype 6): factor wall-clock (left) and peak memory (right) vs. problem size: RSLAB vs. faer vs. MKL PARDISO, with power-law fits.




Figure 3: LU path (unsymmetric, PARDISO mtype 13): factor wall-clock (left) and peak memory (right) vs. problem size. The fixed RSLAB default configuration (gray) accompanies the heuristic pick (blue); on time the pick scales flatter than PARDISO ($\alpha \approx 1.01$ vs. 1.25).

RSLAB vs	$\text{LDL}^\top$ (sym)	LU (unsym)
MKL PARDISO: factor time	$5.6\times$ slower	$5.1\times$ slower
MKL PARDISO: peak memory	$2.8\times$ more	$3.6\times$ more
faer $\text{LU}^\dagger$: factor time	$6.7\times$ faster	$2.7\times$ faster
faer $\text{LU}^\dagger$: peak memory	$1.8\times$ less	$1.4\times$ more
fixed default cfg: factor time	$1.84\times$ faster	$1.49\times$ faster
fixed default cfg: peak memory	$0.90\times$ (less)	$1.11\times$ more

Table 1: Head-to-head geomean ratios (63 sizes per path, 1k–110k DOFs), each path on its own class, over the matrices both solvers factor to below 0.1 residual. The last rows are the shipped heuristic pick against the fixed RSLAB default configuration. Unsym memory ratios above 1 are the worker-count trade: the calibrated pick runs more workers than the capped fixed config, raising the transient panel peak. † faer has no symmetric path (it factors symmetric matrices as LU), so its $\text{LDL}^\top$ gap is structurally largest; it also OOMs on the largest matrices, so its head-to-head is a conservative floor.




Figure 4: Relative residual $\|Ax - b\|/\|b\|$ across the corpus, per solver.

9.2 Apple Silicon: RSLAB vs. Accelerate sparse solvers

A second, independent measurement of all three solver paths on an Apple M3 (8 cores, 16 GB, macOS 15.6) against Apple Accelerate’s sparse direct solvers, which gained complex support and a sparse LU in macOS 15.5. RSLAB runs its shipped default; Accelerate runs its vendor-best configuration per class: Cholesky-first with an $\mathbf{LDL}^\top$ fallback for real-symmetric matrices, its sparse LU (TPP, `pivotTolerance` 0.01) for unsymmetric and circuit-shaped ones, and, because its complex $\mathbf{LDL}^\top$ is Hermitian-only (verified empirically: the lower triangle is mirrored conjugated), LU on the *full* matrix for the complex-symmetric families, the same structural handicap faer carries. Each matrix gets an AMD-vs-Metis ordering bakeoff (`SparseOrderDefault` is plain AMD, up to $4.4\times$ the Metis fill on the large EM/FEM systems); the bakeoff cost is charged to Accelerate’s analyze time, mirroring RSLAB’s exact ordering race.

Results are reported as *wall time divided by Accelerate’s*, so 1.0 is Accelerate and lower is faster, under two metrics that answer different questions: *factor only* is the cost of a repeated factorization on a fixed pattern, *one-shot* is analyze plus factor plus solve, what a caller solving a single system waits for. Both solvers race orderings inside their analyze, so the one-shot column compares like with like. RSLAB’s shipped thread policy caps at four workers while Accelerate uses all cores; on the convection–diffusion class that cap costs 12 to 17%. The numbers below keep the shipped default.

The vendor’s AMX-backed kernels own the small and mid sizes, and the ratio improves with the problem: Helmholtz reaches 0.52 at $n = 110\,592$ and convection–diffusion 3D 0.68 at the same




Figure 5: Apple M3, 8 threads: RSLAB v0.29 wall time divided by Accelerate’s, per matrix class, factor only and one-shot. Bars below 1.0 are faster than the vendor library.

matrix class	factor only	one-shot
circuit MNA (KLU path)	0.42	0.18
Stokes saddle point	0.68	0.43
Helmholtz 3D	0.73	0.47
curl-curl Maxwell	0.74	0.61
convection-diffusion 3D	1.04	0.64
convection-diffusion 2D	1.06	0.35
BEM/MoM near field	1.26	1.49
geomean over the three paths	0.70	0.38

Table 2: Apple M3 head-to-head, RSLAB wall time / Accelerate wall time (lower is faster), over the matrices both solve to below 0.1 residual. Nine sizes per generator family on a 4k–110k log grid plus eight circuit sizes to 200k.

size (Fig. 6). The circuit class stays below 0.45 at every measured size. The one class behind on both metrics is the near-dense BEM/MoM block, which is also the class where the analyze does not pay for itself. Its fronts are medium-sized and near-dense, the regime where the packed complex GEMM runs at a few GF/s against the 31 to 40 GF/s it sustains on hot operands of the same shape, because each update streams its operands once and reuses nothing.

Measurement protocol: absolute times on this machine drift with its thermal state over a long session, so each configuration pair is measured interleaved and reported as the minimum over rounds, and Accelerate is re-measured in the same rounds as the solver it is compared against.

On accuracy, Accelerate reaches 10^{-8} on 25 of 36 corpus matrices, but on five hard indefinite/ill-scaled ones (stokes64, bratu3d, cont-201, rim, lhr34) it returns NaN or garbage *with an OK status*, the same silent-degradation mode as faer, where RSLAB declines cleanly.

9.3 KLU path on circuit-shaped matrices

Table 3 benchmarks the KLU path against the multifrontal LU (defaults) on MNA-like matrices: ~ 4 – 5 nnz/column, unsymmetric, column-diagonally dominant, cascaded stages giving a genuinely reducible BTF structure (`cargo bench --bench klu_circuit`, Apple M3). The KLU factor is 5 to $12\times$ faster with 1.7 to $5.7\times$ less fill, the gap widening with size as the multifrontal fronts grow; the




Figure 6: Normalized wall time vs. problem size, one line per matrix class. Every class improves with size; the two that end above 1.0 are the ones whose fronts stay small (convection–diffusion) or near-dense (BEM/MoM).

numeric-only refactor runs $\sim 2\times$ faster still, so a 20-point sweep (refactor+solve vs. factor+solve) is 10 to $40\times$ faster end to end. Both solvers reach machine-precision residuals ($\sim 10^{-15}$) on every size. The shipped default runs the independent BTF blocks in parallel once a deterministic structural gate clears (at least 4 blocks, 8k nonzeros), bit-identical to the sequential result, adding 3.1 to $3.9\times$ on factor and 3.9 to $4.5\times$ on refactor.

n	nnz	KLU				multifrontal LU		
		factor	refactor	fill	blocks	factor	fill	sweep ratio
2k	15k	1.7 ms	0.8 ms	79k	8	8.1 ms	132k	9.8 $\times$
10k	73k	8.0 ms	4.1 ms	439k	16	54.7 ms	1.03M	12.5 $\times$
50k	366k	45.4 ms	24.3 ms	2.32M	32	304 ms	9.21M	13.9 $\times$
200k	1.47M	179 ms	96 ms	9.17M	64	2.16 s	52.2M	40 $\times$

Table 3: KLU vs. multifrontal **LU** on MNA-like circuit matrices (Apple M3). “Sweep ratio” is the wall-clock ratio of 20 same-pattern value sets solved by KLU refactor+solve vs. multifrontal factor+solve. Fill counts stored factor entries.

Against the reference C implementation, SuiteSparse KLU [14] (loaded at runtime when installed), the analyze pipeline is at parity: identical BTF block counts and fill within 1.5% (RSLAB slightly less), confirming the maximum transversal, Tarjan SCC, and per-block AMD match the reference. On the numeric kernel (32-bit index streams, packed DFS marks, Eisenstat-Liu symmetric pruning [13], and a recorded refactor scatter program), the sequential factor runs within 1.2 to $1.6\times$ of the C code (e.g. 170 ms vs. 135 ms at $n = 200k$) and the refactor within 1.15 to $1.4\times$. The bit-identical parallel per-block execution (the gated shipped default; factor and refactor; two-phase spliced output) is 1.6 to $2.7\times$ faster than SuiteSparse KLU on factor and 2.1 to $3.4\times$ on refactor (51 ms / 24 ms at $n = 200k$), putting the 20-point sweep $\sim 2.4\times$ ahead end to end. Both reach $\sim 10^{-15}$ residuals. On the real SuiteSparse circuit corpus the column pipeline inside the dominant block (Sec. 4.5) adds 2.0 to $2.9\times$ on refactor on the work-heavy matrices, where the block structure is one large irreducible block that per-block parallelism cannot split.




Figure 7: Real circuit matrices from the SuiteSparse collection: RSLAB KLU vs. SuiteSparse KLU, one pivoting factorization (left) and a 20-point same-pattern sweep of refactor plus solve (right), log scale, labelled with the per-matrix ratio. The ratio is matrix-dependent, spanning $0.47\times$ to $27.9\times$ on the factor and $0.38\times$ to $25.8\times$ on the sweep.




Figure 8: Strong scaling of the block solve against one worker, per block width, with the per-RHS MGS path (dashed) for reference.

9.4 Iterative layer

The Krylov results appear with their concepts in Sec. 8: block-CGS2 lifts multi-RHS strong scaling to $\sim 2.2\times$ at 12 cores where per-RHS MGS is flat; within-cycle deflation cuts the batched operator column-applies to $0.66\times$ the full-width bound at $s = 16$; GCRO-DR recycling cuts the cross-solve iteration total $6.4\times$ on a stagnating sequence; and the incomplete-factor sweet spot at $\text{drop_tol} = 10^{-2}$ halves the factor memory at a total wall time within a few percent of the exact direct solve. FGMRES’s flexible-basis update saves exactly one preconditioner solve per restart cycle, confirmed by counting.




Figure 9: A sequence of eight related solves: cumulative wall-clock (left) and cumulative GMRES iterations (right), cold start vs. warm start vs. GCRO-DR recycling at three subspace sizes.

10 Conclusion

RSLAB is a pure-Rust, scalar-generic sparse direct solver with three paths matched to their operator classes: a supernodal complex-symmetric $\mathbf{LDL}^T$, a supernodal/multifrontal threshold-pivoted $\mathbf{LU}$, and a sequential, bit-deterministic KLU path with a numeric-only refactor for fixed-pattern sweeps. It reaches within a single-digit factor of MKL PARDISO while remaining dependency-free, is faster and leaner than the open-source alternatives, and on Apple Silicon it factors the corpus in 0.70 of Accelerate’s wall time (0.38 one-shot) in geomean over the three paths. Its analyze phase computes memory, work and runtime as exact functions of the structure, which is what lets the shipped configuration be an exact race over candidates rather than a model of them. When the factor is made approximate it becomes a preconditioner, and the matching Krylov layer (COCG/COCR, flexible GMRES, deflating block GMRES, GCRO-DR recycling) turns the direct factor into the engine of a fast preconditioned iterative solver.

References

- [1] P. R. Amestoy, T. A. Davis, I. S. Duff. An approximate minimum degree ordering algorithm. *SIAM J. Matrix Anal. Appl.* 17(4):886–905, 1996.
- [2] P. R. Amestoy. Recent progress in parallel multifrontal solvers and the approximate minimum fill (HAMF) ordering. Technical report / CERFACS, 1999.
- [3] E. Rothberg, S. C. Eisenstat. Node selection strategies for bottom-up sparse matrix ordering. *SIAM J. Matrix Anal. Appl.* 19(3):682–695, 1998.
- [4] A. George. Nested dissection of a regular finite element mesh. *SIAM J. Numer. Anal.* 10(2):345–363, 1973.
- [5] G. Karypis, V. Kumar. A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM J. Sci. Comput.* 20(1):359–392, 1998.
- [6] J. D. Hogg, E. Ovtchinnikov, J. A. Scott. A sparse symmetric indefinite direct solver for GPU architectures (SSIDS). *ACM Trans. Math. Softw.* 42(1):1–25, 2016.
- [7] C. Ashcraft, R. Grimes. The influence of relaxed supernode partitions on the multifrontal method. *ACM Trans. Math. Softw.* 15(4):291–309, 1989.
- [8] J. R. Bunch, L. Kaufman. Some stable methods for calculating inertia and solving symmetric linear systems. *Math. Comp.* 31(137):163–179, 1977.
- [9] I. S. Duff, J. K. Reid. The multifrontal solution of indefinite sparse symmetric linear equations. *ACM Trans. Math. Softw.* 9(3):302–325, 1983.
- [10] J. W. H. Liu. The multifrontal method for sparse matrix solution: theory and practice. *SIAM Review* 34(1):82–109, 1992.

- [11] J. W. H. Liu. On the storage requirement in the out-of-core multifrontal method for sparse factorization. *ACM Trans. Math. Softw.* 12(3):249–264, 1986.
- [12] T. A. Davis. Algorithm 832: UMFPACK, an unsymmetric-pattern multifrontal method. *ACM Trans. Math. Softw.* 30(2):196–199, 2004.
- [13] S. C. Eisenstat, J. W. H. Liu. Exploiting structural symmetry in unsymmetric sparse symbolic factorization. *SIAM J. Matrix Anal. Appl.*, 13(1), 1992.
- [14] T. A. Davis, E. Palamadai Natarajan. Algorithm 907: KLU, a direct sparse solver for circuit simulation problems. *ACM Trans. Math. Softw.* 37(3):36, 2010.
- [15] E. Carson, N. J. Higham. Accelerating the solution of linear systems by iterative refinement in three precisions. *SIAM J. Sci. Comput.* 40(2):A817–A847, 2018.
- [16] J. E. Hopcroft, R. M. Karp. An $n^{5/2}$ algorithm for maximum matchings in bipartite graphs. *SIAM J. Comput.* 2(4):225–231, 1973.
- [17] X. Chen, Y. Wang, H. Yang. NICSLU: an adaptive sparse matrix solver for parallel circuit simulation. *IEEE Trans. CAD* 32(2):261–274, 2013.
- [18] O. Schenk, K. Gärtner. Solving unsymmetric sparse systems of linear equations with PARDISO. *Future Gener. Comput. Syst.* 20(3):475–487, 2004.
- [19] S. Quinones. faer-rs: a linear algebra library for the Rust programming language. <https://github.com/sarah-quinones/faer-rs>, 2024.
- [20] T. A. Davis, Y. Hu. The University of Florida sparse matrix collection. *ACM Trans. Math. Softw.* 38(1):1–25, 2011.
- [21] Y. Saad, M. H. Schultz. GMRES: a generalized minimal residual algorithm for solving nonsymmetric linear systems. *SIAM J. Sci. Stat. Comput.* 7(3):856–869, 1986.
- [22] Y. Saad. A flexible inner–outer preconditioned GMRES algorithm. *SIAM J. Sci. Comput.* 14(2):461–469, 1993.
- [23] J. W. Daniel, W. B. Gragg, L. Kaufman, G. W. Stewart. Reorthogonalization and stable algorithms for updating the Gram–Schmidt QR factorization. *Math. Comp.* 30(136):772–795, 1976.
- [24] L. Giraud, J. Langou, M. Rozložník, J. van den Eshof. Rounding error analysis of the classical Gram–Schmidt orthogonalization process. *Numer. Math.* 101(1):87–100, 2005.
- [25] J. L. Barlow, A. Smoktunowicz. Reorthogonalized block classical Gram–Schmidt. *Numer. Math.* 123(3):395–423, 2013.
- [26] M. L. Parks, E. de Sturler, G. Mackey, D. D. Johnson, S. Maiti. Recycling Krylov subspaces for sequences of linear systems. *SIAM J. Sci. Comput.* 28(5):1651–1674, 2006.
- [27] R. B. Morgan. GMRES with deflated restarting. *SIAM J. Sci. Comput.* 24(1):20–37, 2002.
- [28] H. A. van der Vorst, J. B. M. Melissen. A Petrov–Galerkin type method for solving $Ax = b$, where A is symmetric complex. *IEEE Trans. Magn.* 26(2):706–708, 1990.
- [29] T. Sogabe, S.-L. Zhang. A COCR method for solving complex symmetric linear systems. *J. Comput. Appl. Math.* 199(2):297–303, 2007.